



Mahidol University *Wisdom of the Land*

Chapter 9

Image Compression

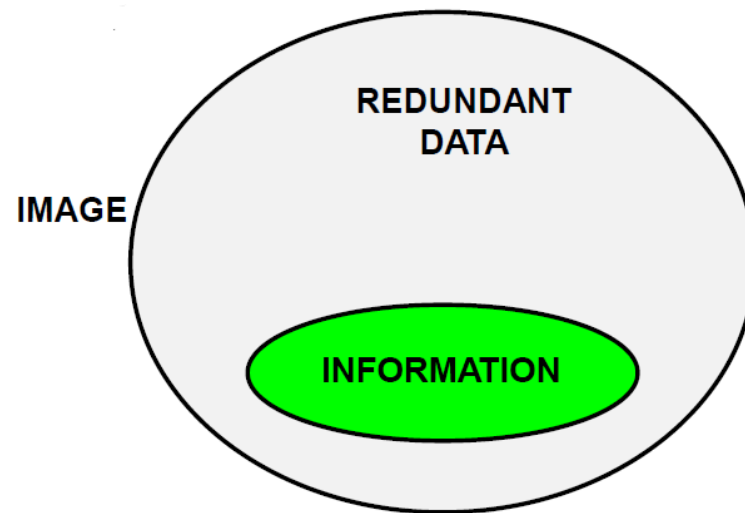
Compression Fundamentals

- Data compression
 - the process of reducing the amount of data required to represent a given quantity of information.
- Data and Information
 - Note that data and information are not the same.
 - Data are the means by which information is conveyed.
 - Various amounts of data may be used to represent the same amount of information.
- Data that either provide no relevant information or simply restate that which is already known
 - data redundancy

Compression Fundamentals

- **Data vs Information**

- Information = Matter
- Data = The means by which information is conveyed

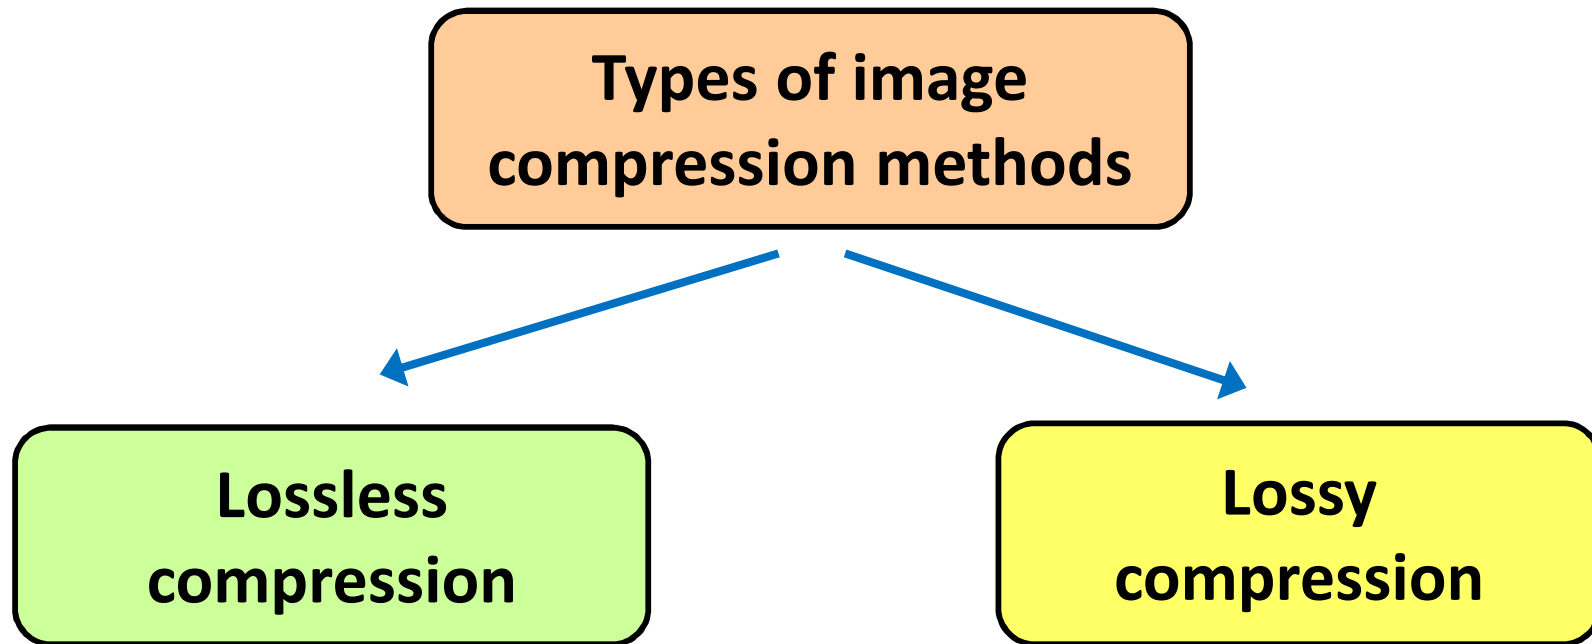


$$\text{IMAGE} = \text{INFORMATION} + \text{REDUNDANT DATA}$$

- **Image Compression**

- Reducing the amount of data required to represent a digital image while keeping information as much as possible

Image Compression Methods



Lossless Compression VS. Lossy Compression

Lossless compression

- Information preserving
 - Low compression ratios
 - Zero data loss
-
- Example: Fixed length coding and Variable length coding

Lossy compression

- Not information preserving
 - High compression ratios
 - Some data loss
-
- Example: JPEG compression, JPEG2000 compression.

- **Coding Redundancy** : A code is a system of symbols (i.e. bits) that represents information. Each piece of information is represented by a set of code symbols.
- The normalized histogram of a gray level image can be used in construction of codes to reduce the data used to represent it.

$$p_r(r_k) = \frac{n_k}{n}, k = 0, 1, 2, \dots, L-1$$

where

r_k is the pixel values defined in the interval $[0,1]$

$p_r(r_k)$ is the probability of occurrence of r_k .

L is the number of gray levels.

n_k is the number of times that k^{th} gray level appears

n is the total number of pixels.

- Different coding methods yield different amount of data needed to represent the same information.
- Average number of bits required to represent each pixel is given by :

$$L_{avg} = \sum_{k=0}^{L-1} l(r_k) p_r(r_k)$$

where

$l(r_k)$ is the number of bits used to represent each pixel of r_k .

Relative Data Redundancy and Compression Ratio

- Given n_1 and n_2 denoting the information-carrying units in two data sets that represent the same information/image.
- Relative data redundancy, R_D of the first data set, n_1 , is defined by :

$$R_D = 1 - \frac{1}{C_R}$$

- C_R refers to the compression ratio and is defined by :

$$C_R = \frac{n_1}{n_2}$$

- If $n_1 = n_2$, then $C_R = 1$ and $R_D = 0$, indicating that the first representation of the information contains no redundant data.
- A typical compression ratio around 15 or (15:1) indicates that 85% of the data in the first data set is redundant.

Example : Fixed Length Coding

- Assuming there is an image 3 bit ($L=8$). There are gray levels distribution values as Table in which the gray levels range $[0, L-1] = [0, 7]$.

r_k	$p_r(r_k)$	Fixed Length Coding	$I_r(r_k)$
$r_0 = 0$	0.19	000	3
$r_1 = 1/7$	0.25	001	3
$r_2 = 2/7$	0.21	010	3
$r_3 = 3/7$	0.16	011	3
$r_4 = 4/7$	0.08	100	3
$r_5 = 5/7$	0.06	101	3
$r_6 = 6/7$	0.03	110	3
$r_7 = 1$	0.02	111	3

Example : Fixed Length Coding

- The average number of bit used for fixed 3-bit code :

$$\begin{aligned} L_{avg} &= \sum_{k=0}^{L-1} l(r_k) p_r(r_k) \\ &= 3(0.19) + 3(0.25) + 3(0.21) + 3(0.16) + 3(0.08) + 3(0.06) + 3(0.03) + 3(0.02) \\ &= 3.1 \\ &= 3 \text{ bits} \end{aligned}$$

Example : Variable Length Coding

- Assuming there is an image 3 bit ($L=8$). There are gray levels distribution values as Table in which the gray levels range $[0, L-1] = [0, 7]$.

r_k	$p_r(r_k)$	Variable Length Coding	$l_r(r_k)$
$r_0 = 0$	0.19	11	2
$r_1 = 1/7$	0.25	01	2
$r_2 = 2/7$	0.21	10	2
$r_3 = 3/7$	0.16	001	3
$r_4 = 4/7$	0.08	0001	4
$r_5 = 5/7$	0.06	00001	5
$r_6 = 6/7$	0.03	000001	6
$r_7 = 1$	0.02	000000	6

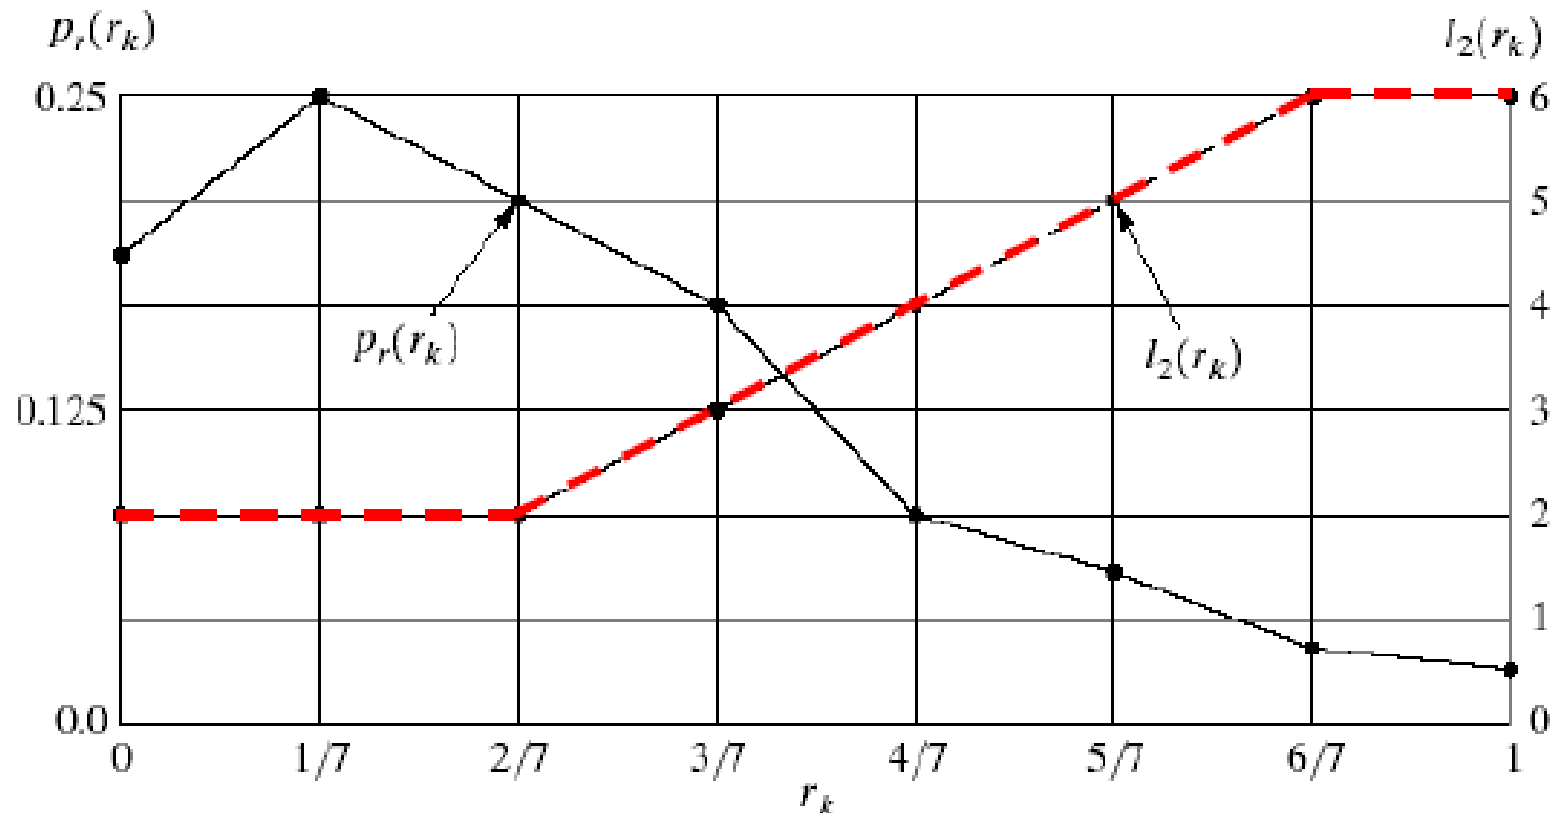
Example : Variable Length Coding

- The average number of bit used for variable length code :

$$\begin{aligned} L_{avg} &= \sum_{k=0}^{L-1} l(r_k) p_r(r_k) \\ &= 2(0.19) + 2(0.25) + 2(0.21) + 3(0.16) + 4(0.08) + 5(0.06) + 6(0.03) + 6(0.02) \\ &= 2.7 \text{ bits} \end{aligned}$$

Example : Variable Length Coding

- Data compression is achieved by assigning fewer bits to more probable gray levels than the less probable gray levels.



Concept : assign the longest code word to the symbol with the least probability of occurrence.

Fixed Length Coding VS. Variable Length Coding

- Assuming there is an image 3 bit ($L=8$). There are gray levels distribution values as Table in which the gray levels range $[0, L-1] = [0, 7]$.

r_k	$p_r(r_k)$	Code 1	$l_1(r_k)$	Code 2	$l_2(r_k)$
$r_0 = 0$	0.19	000	3	11	2
$r_1 = 1/7$	0.25	001	3	01	2
$r_2 = 2/7$	0.21	010	3	10	2
$r_3 = 3/7$	0.16	011	3	001	3
$r_4 = 4/7$	0.08	100	3	0001	4
$r_5 = 5/7$	0.06	101	3	00001	5
$r_6 = 6/7$	0.03	110	3	000001	6
$r_7 = 1$	0.02	111	3	000000	6

Fixed 3-bit code

Variable length code

Fixed Length Coding VS. Variable Length Coding

- The average number of bit used for fixed 3-bit code :

$$\begin{aligned} L_{avg} &= \sum_{k=0}^{L-1} l(r_k) p_r(r_k) \\ &= 3(0.19) + 3(0.25) + 3(0.21) + 3(0.16) + 3(0.08) + 3(0.06) + 3(0.03) + 3(0.02) \\ &= 3 \text{ bits} \end{aligned}$$

- The average number of bit used for variable length code :

$$\begin{aligned} L_{avg} &= \sum_{k=0}^{L-1} l(r_k) p_r(r_k) \\ &= 2(0.19) + 2(0.25) + 2(0.21) + 3(0.16) + 4(0.08) + 5(0.06) + 6(0.03) + 6(0.02) \\ &= 2.7 \text{ bits} \end{aligned}$$

Fixed Length Coding VS. Variable Length Coding

- The compression ratio :

$$C_R = \frac{n_1}{n_2} \begin{array}{l} \leftarrow \text{Fixed 3-bit code} \\ \leftarrow \text{Variable length code} \end{array}$$
$$= \frac{3}{2.7} = 1.11$$

- The relative Data Redundancy :

$$R_D = 1 - \frac{1}{C_R}$$
$$= 1 - \frac{1}{1.11} = 0.099 \Rightarrow \cong \%10$$

Elements of Information Theory

- **Measuring Information** : The information in an image can be modeled as a probabilistic process, where we first develop a statistical model of the image generation process. The information content (entropy) can be estimated based on this model.
- The information per source (symbol or pixel), which is also referred as entropy is calculated by :

$$e = -\sum_{j=1}^J P(a_j) \log_2 P(a_j)$$

- Where $P(a_j)$ refers to the source symbol/pixel probabilities. J refers to the number of symbols or different pixel values.

Measuring Information

- For example, given the following gray scale image :

16	16	16	90	164	238	238	238
16	16	16	90	164	238	238	238
16	16	16	90	164	238	238	238
16	16	16	90	164	238	238	238

Measuring Information

- The entropy of the given gray scale mage can be calculated by :

Gray Level	Count	Probability
16	12	3/8
90	4	1/8
161	4	1/8
238	12	3/8

- The entropy of this image is calculated by :

$$\begin{aligned} e &= -\sum_{j=1}^J P(a_j) \log_2 P(a_j) \\ &= -[(3/8) \log_2 (3/8) + (1/8) \log_2 (1/8) + (1/8) \log_2 (1/8) + (3/8) \log_2 (3/8)] \\ &= 1.81 \text{ bits / pixel} \end{aligned}$$

- **Huffman Coding** : The Huffman coding involves the following 2 steps.
 - 1) Create a series of source reductions by ordering the probabilities of the symbols and combining the lowest probability symbols into a single symbol and replace in the next source reduction.
 - 2) Code each reduced source starting with the smallest source and working back to the original source. Use 0 and 1 to code the simplest 2 symbol source.

Huffman Coding

- 1) Huffman source reductions : a_i 's corresponds to the available gray levels in a given image.

Original source		Source reduction			
Symbol	Probability	1	2	3	4
a_2	0.4	0.4	0.4	0.4	0.6
a_6	0.3	0.3	0.3	0.3	
a_1	0.1	0.1	0.2	0.3	0.4
a_4	0.1	0.1			
a_3	0.06	0.1			
a_5	0.04				

Huffman Coding

- 2) Huffman code assignments : The first code assignment is done for a_2 with the highest probability and the last assignments are done for a_3 and a_5 with the lowest probabilities.

Original source			Source reduction							
Sym.	Prob.	Code	1		2		3		4	
a_2	0.4	1	0.4	1	0.4	1	0.4	1	0.6	0
a_6	0.3	00	0.3	00	0.3	00	0.3	00	0.4	1
a_1	0.1	011	0.1	011	0.2	010	0.3	01		
a_4	0.1	0100	0.1	0100	0.1	011				
a_3	0.06	01010	0.1	0101						
a_5	0.04	01011								

Last code

First code

The code is instantaneous uniquely decodable without referencing succeeding symbols.

Huffman Coding

- Note that the shortest codeword (1) is given for the symbol/pixel with the highest probability (a_2). The longest codeword (01011) is given for the symbol/pixel with the lowest probability (a_5).
- The average length of the Huffman code is given by :

$$\begin{aligned} L_{avg} &= (0.4)(1) + (0.3)(2) + (0.1)(3) + (0.1)(4) + (0.06)(5) + (0.04)(5) \\ &= 2.2 \text{ bits / symbol} \end{aligned}$$

- The entropy of the Huffman code is given by :

$$\begin{aligned} e &= -\sum_{j=1}^J P(a_j) \log_2 P(a_j) \\ &= 2.14 \text{ bits / symbol} \end{aligned}$$

- Huffman code efficiency is given by :

$$C = \frac{e}{L_{avg}} \times 100$$

where

L_{avg} is the average length of the Huffman code

e is the entropy of the Huffman code

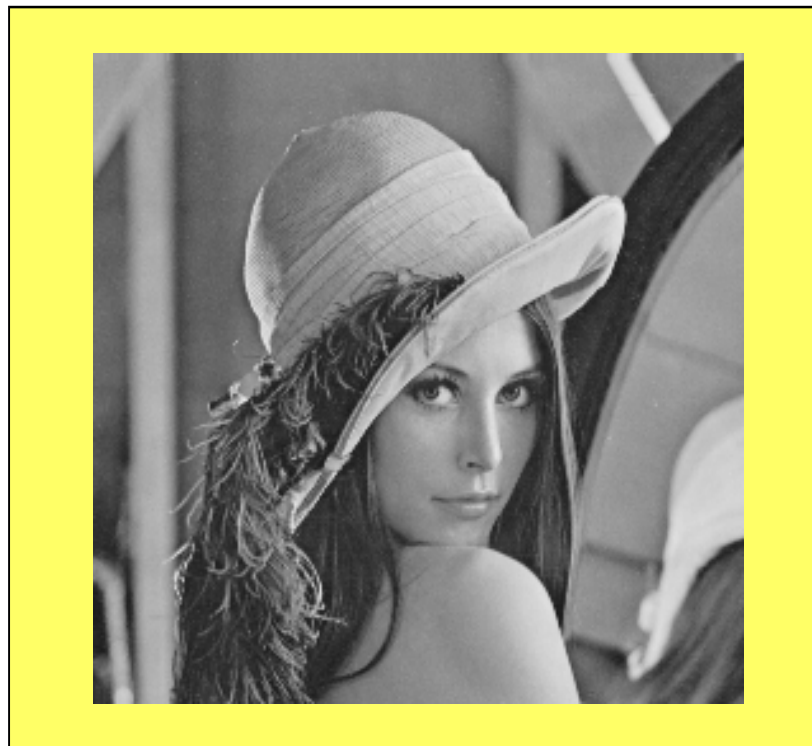
- The resulting Huffman coding efficiency is 97.3% $((2.14/2.2)*100)$.

JPEG compression

- Story of JPEG
 - Joint Photographic Expert Group
- International Organization for Standards (ISO)
 - 1988: ISO got together a group of experts to develop a good image compression scheme
 - Geared towards photographs of natural imagery
 - Color and monochrome
 - Easy to use (spin-dial quality control)
 - Through empirical testing, the following scheme proved to be the best
 - Standardized in August 1990

Block partition

- First, the image is divided into 8×8 non-overlapping blocks, which are processed left to right, top to bottom.



Image

Partition
→

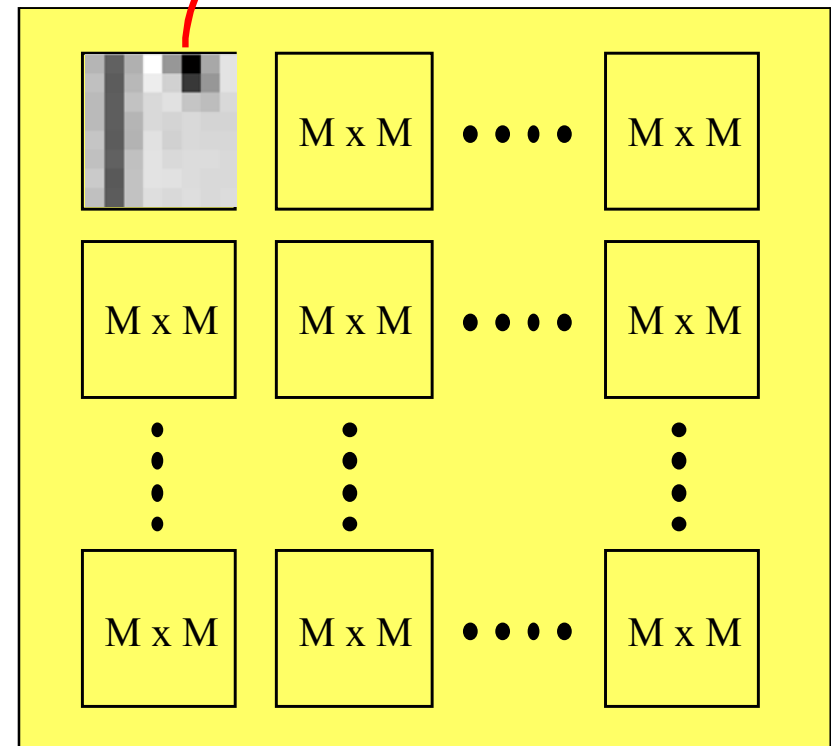
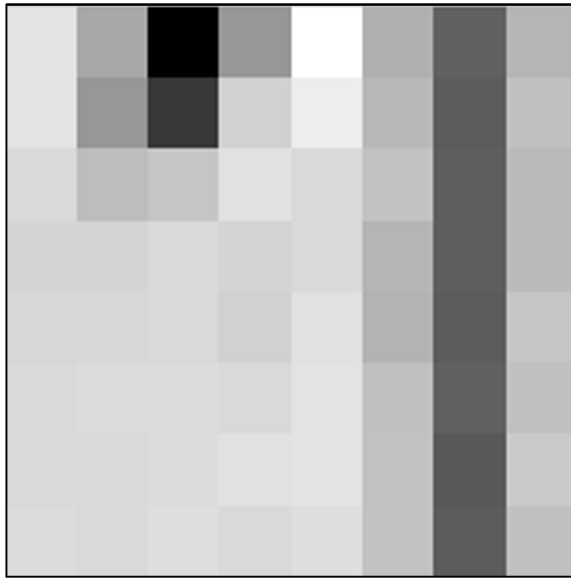


Image blocks

Example of image block

- As an example, 8×8 block of 8-bit image might be:

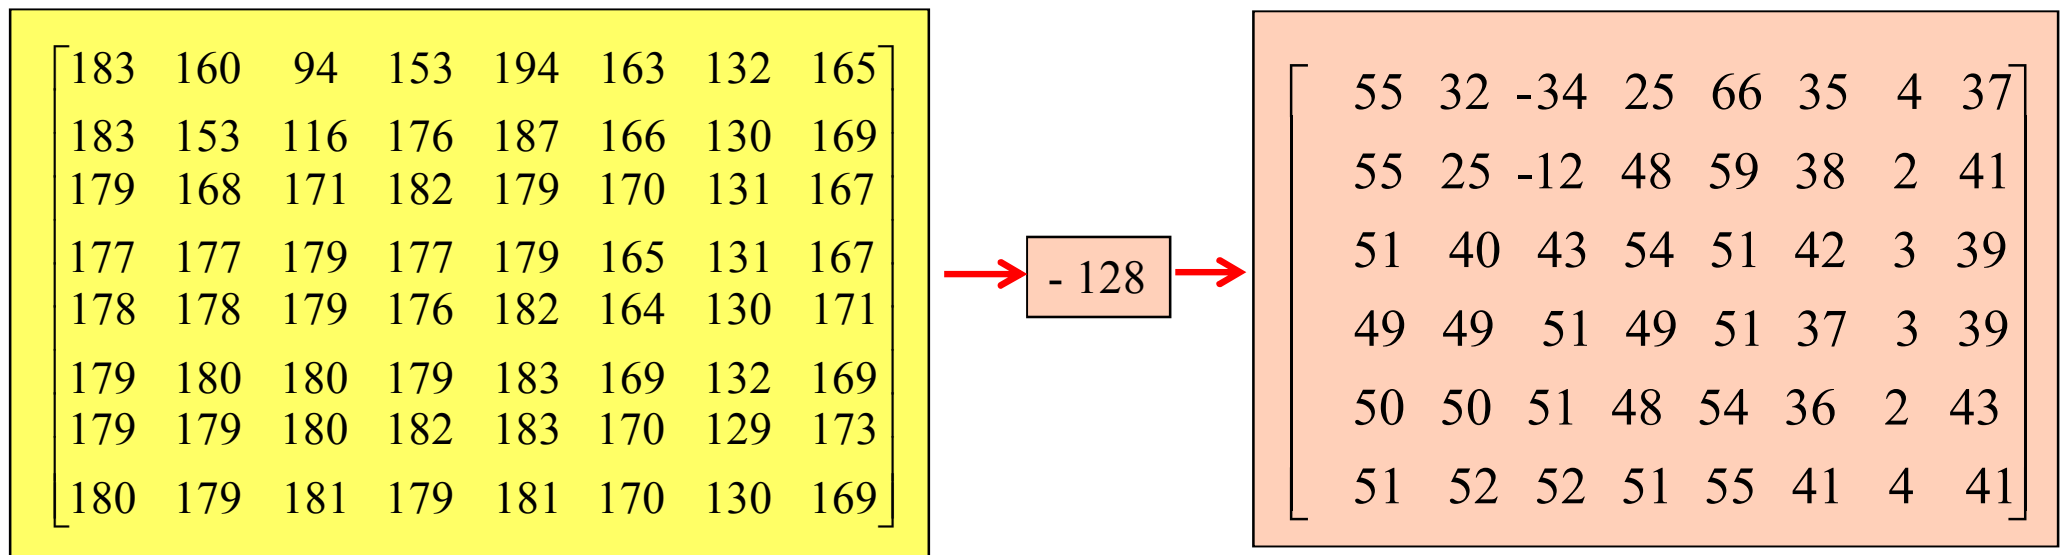


183	160	94	153	194	163	132	165
183	153	116	176	187	166	130	169
179	168	171	182	179	170	131	167
177	177	179	177	179	165	131	167
178	178	179	176	182	164	130	171
179	180	180	179	183	169	132	169
179	179	180	182	183	170	129	173
180	179	181	179	181	170	130	169

Image block

Shift value procedure

- Before computing the DCT of the image block, its gray values are shifted from a positive range to one centered around zero.
- For an 8-bit image each pixel has 256 possible values: $[0, 255]$. To center around zero it is necessary to subtract by half the number of possible values, or 128.



- As an example, the first number is derived from $183 - 128 = 55$.

Discrete Cosine Transform (DCT)

- 2D-DCT:
$$G_{u,v} = \sum_{x=0}^7 \sum_{y=0}^7 \alpha(u)\alpha(v)g_{x,y} \cos\left[\frac{\pi}{8}\left(x + \frac{1}{2}\right)u\right] \cos\left[\frac{\pi}{8}\left(y + \frac{1}{2}\right)v\right]$$

Image block

55	32	-34	25	66	35	4	37
55	25	-12	48	59	38	2	41
51	40	43	54	51	42	3	39
49	49	51	49	51	37	3	39
50	50	51	48	54	36	2	43
51	52	52	51	55	41	4	41

DCT

DC coefficient

low frequency → high frequency

low frequency

high frequency

313.0	55.8	-27.1	17.8	78.0	-59.8	26.7	-26.6
-38.2	-27.2	13.2	44.3	31.7	-0.9	-23.8	-9.9
-20.1	-17.4	9.8	32.8	21.2	-6.1	-15.8	-8.6
-10.3	-8.0	8.7	16.5	8.8	-10.2	-12.9	0.5
-6.3	1.2	6.4	4.0	-3.3	-7.5	-5.4	5.3
2.5	3.3	0.4	-2.6	-6.8	-3.8	0.5	2.5
3.7	4.4	-1.1	-2.0	-9.0	-0.2	2.4	3.5
2.5	1.1	-0.4	-3.6	-1.6	-1.0	3.1	1.5

DCT block

29

Quantization table

- The human eye is good at seeing small differences in brightness over a relatively large area, but not so good at distinguishing the exact strength of a high frequency brightness variation. This allows one to greatly reduce the amount of information in the high frequency components.
- The JPEG recommended luminance and chrominance quantization tables can be scaled to provide a variety of compression levels (select the quality of JPEG compression).

Luminance

16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	57	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101
72	92	95	98	112	100	103	99

Chrominance

17	18	24	47	99	99	99	99
18	21	26	66	99	99	99	99
24	26	56	99	99	99	99	99
47	66	99	99	99	99	99	99
99	99	99	99	99	99	99	99
99	99	99	99	99	99	99	99
99	99	99	99	99	99	99	99
99	99	99	99	99	99	99	99

Quantization procedure

DCT block

313.0	55.8	-27.1	17.8	78.0	-59.8	26.7	-26.6
-38.2	-27.2	13.2	44.3	31.7	-0.9	-23.8	-9.9
-20.1	-17.4	9.8	32.8	21.2	-6.1	-15.8	-8.6
-10.3	-8.0	8.7	16.5	8.8	-10.2	-12.9	0.5
-6.3	1.2	6.4	4.0	-3.3	-7.5	-5.4	5.3
2.5	3.3	0.4	-2.6	-6.8	-3.8	0.5	2.5
3.7	4.4	-1.1	-2.0	-9.0	-0.2	2.4	3.5
2.5	1.1	-0.4	-3.6	-1.6	-1.0	3.1	1.5



Quantization table

16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	57	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101
72	92	95	98	112	100	103	99

Quantized block

20	5	-3	1	3	-1	1	0
-3	-2	1	2	1	0	0	0
-1	-1	1	1	1	0	0	0
-1	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

$$\text{round}\left(\frac{313}{16}\right) = 20$$

Round

Zigzag procedure

Quantized block

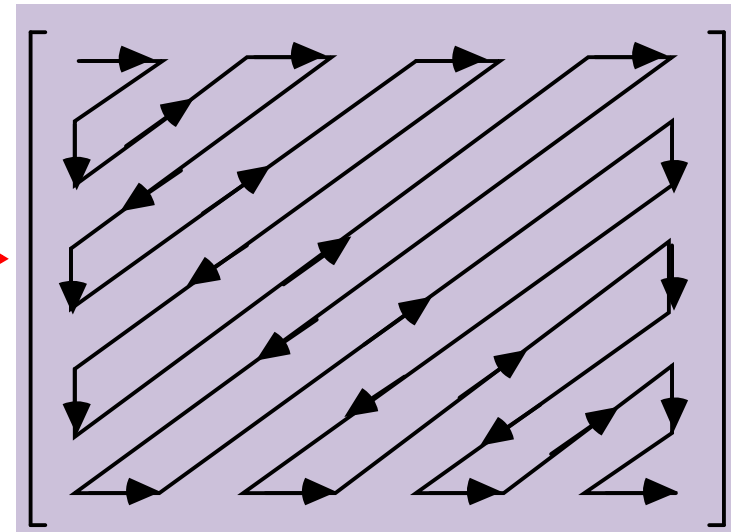
20	5	-3	1	3	-1	1	0
-3	-2	1	2	1	0	0	0
-1	-1	1	1	1	0	0	0
-1	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0



Zigzag
scan



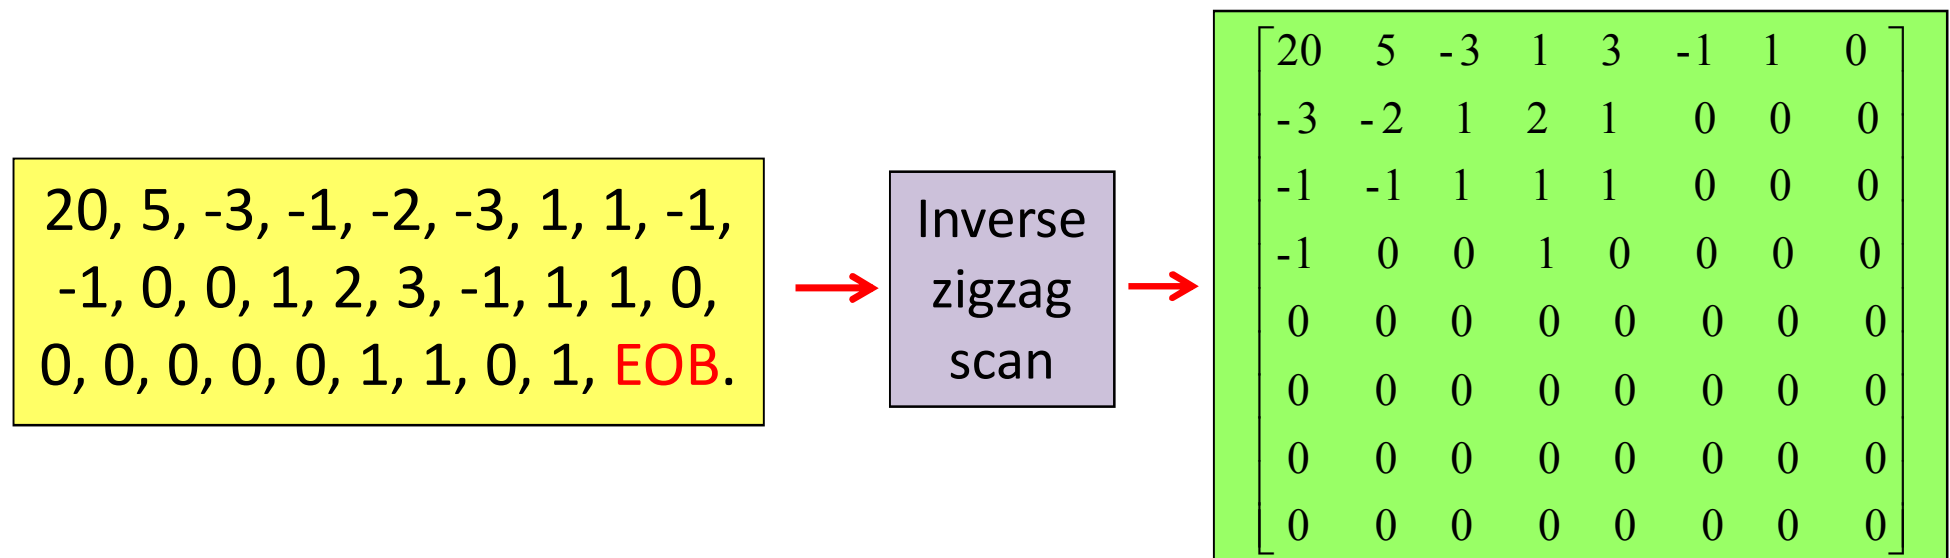
Zigzag scan



- Result = 20, 5, -3, -1, -2, -3, 1, 1, -1, -1, 0, 0, 1, 2, 3, -1, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, **EOB**.
- **EOB** symbol denotes the end-of-block condition.

Inverse zigzag procedure

- Decoding to display the image consists of doing all the above in reverse. Taking the DCT coefficient matrix (after adding the difference of the DC coefficient back in).



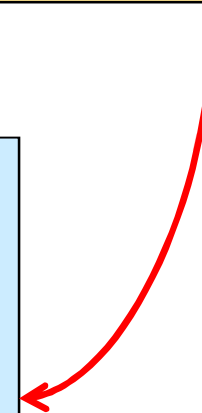
Inverse Quantization procedure

20	5	-3	1	3	-1	1	0
-3	-2	1	2	1	0	0	0
-1	-1	1	1	1	0	0	0
-1	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0



16	11	10	16	24	40	51	61
12	12	14	19	26	58	60	55
14	13	16	24	40	57	69	56
14	17	22	29	51	87	80	62
18	22	37	56	68	109	103	77
24	35	55	64	81	104	113	92
49	64	78	87	103	121	120	101
72	92	95	98	112	100	103	99

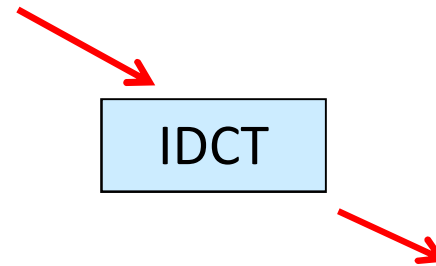
320	55	-30	16	72	-40	51	0
-36	-24	14	38	26	0	0	0
-14	-13	16	24	40	0	0	0
-14	0	0	29	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0



Inverse Discrete Cosine Transform (IDCT)

■ Inverse 2D-DCT:
$$f_{x,y} = \sum_{u=0}^7 \sum_{v=0}^7 \alpha(u)\alpha(v)F_{u,v} \cos\left[\frac{\pi}{8}\left(x + \frac{1}{2}\right)u\right] \cos\left[\frac{\pi}{8}\left(y + \frac{1}{2}\right)v\right]$$

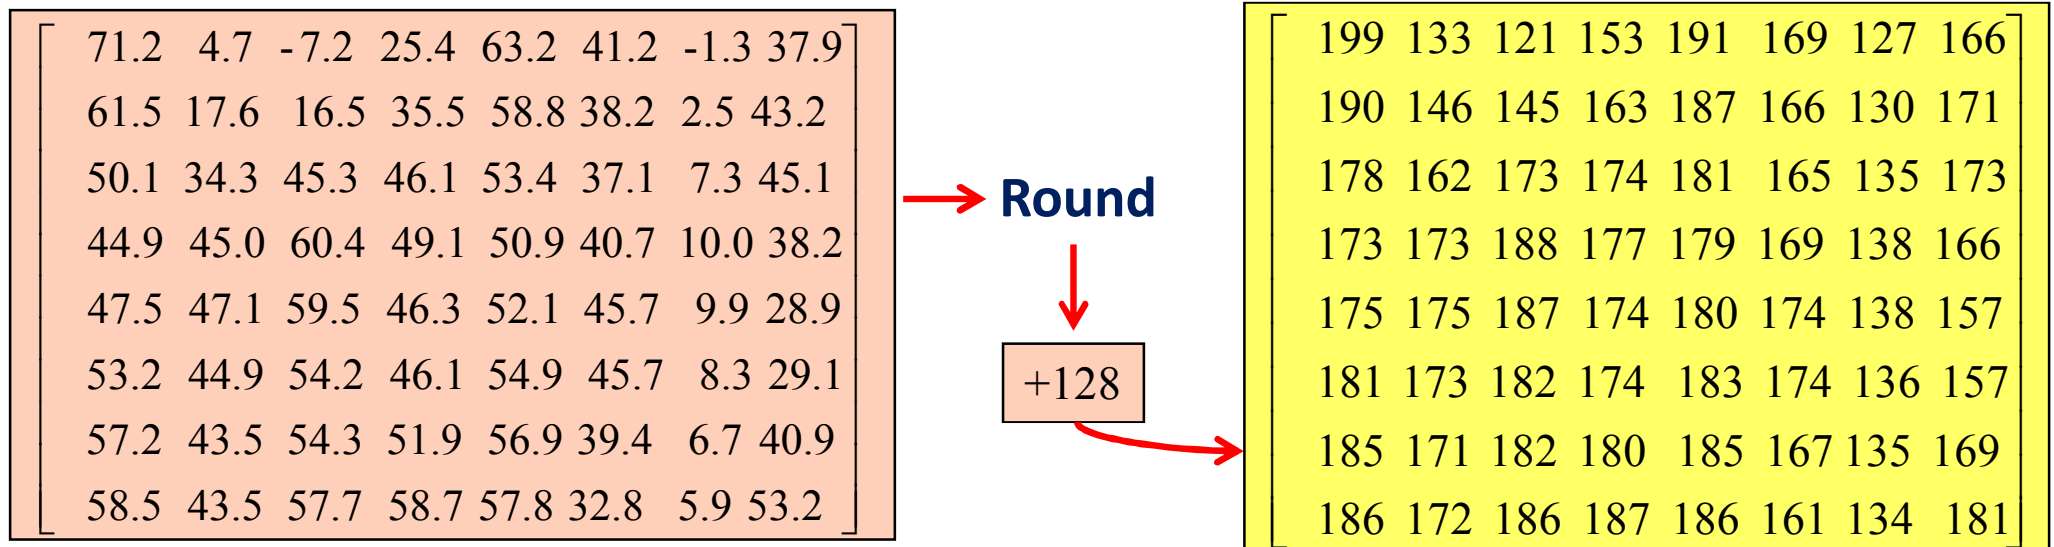
320	55	-30	16	72	-40	51	0
-36	-24	14	38	26	0	0	0
-14	-13	16	24	40	0	0	0
-14	0	0	29	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0



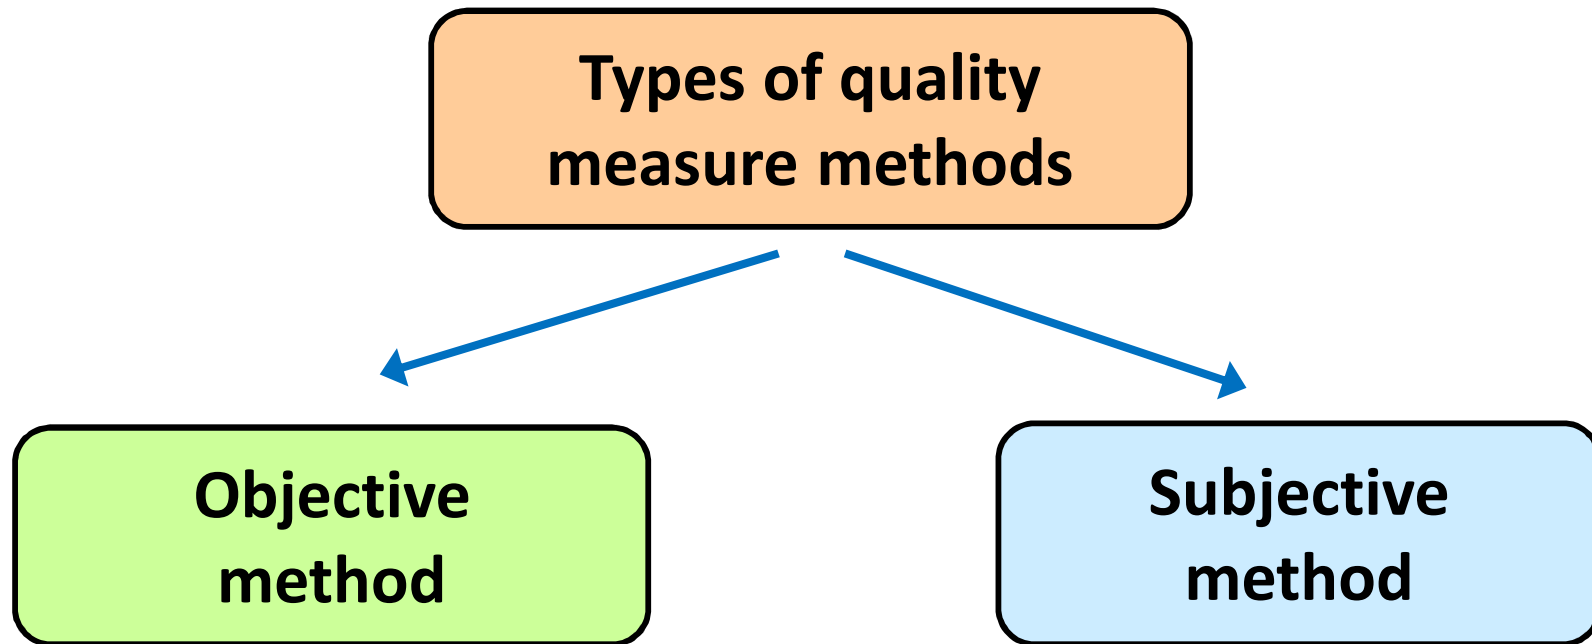
71.2	4.7	-7.2	25.4	63.2	41.2	-1.3	37.9
61.5	17.6	16.5	35.5	58.8	38.2	2.5	43.2
50.1	34.3	45.3	46.1	53.4	37.1	7.3	45.1
44.9	45.0	60.4	49.1	50.9	40.7	10.0	38.2
47.5	47.1	59.5	46.3	52.1	45.7	9.9	28.9
53.2	44.9	54.2	46.1	54.9	45.7	8.3	29.1
57.2	43.5	54.3	51.9	56.9	39.4	6.7	40.9
58.5	43.5	57.7	58.7	57.8	32.8	5.9	53.2

Inverse shift value procedure

- After that, rounding the output to integer values, and adding 128 to each entry. The decompression process may produce values outside of the original input range of [0,255].
- If this occurs, the decoder needs to clip the output values keep them within that range to prevent overflow when storing the decompressed image with the original bit depth.



Quality Measure of a Compressed Image



Quality Measure of a Compressed Image

Objective quality measures (by formulate a criterion) :

- The mean-square error between an original image $\hat{f}(x, y)$ and compressed image $f(x, y)$:

$$e_{MSE} = \frac{1}{MN} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} [\hat{f}(x, y) - f(x, y)]^2$$

where the images are of size $M \times N$.

- Root-mean-square error :

$$e_{RMSE} = \sqrt{\frac{1}{MN} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} [\hat{f}(x, y) - f(x, y)]^2}$$

The smaller the value of e_{RMSE} , the better the compressed image represents the original image.

Quality Measure of a Compressed Image

- Mean-square signal-to-noise ratio :

$$SNR_{ms} = \frac{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} \hat{f}(x, y)^2}{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} [\hat{f}(x, y) - f(x, y)]^2}$$

- Peak signal-to-noise ratio (*PSNR*) – in decibel (dB) :

$$PSNR = 10(\log_{10}) \frac{(L-1)^2}{e_{MSE}}$$

where L is the number of gray levels.

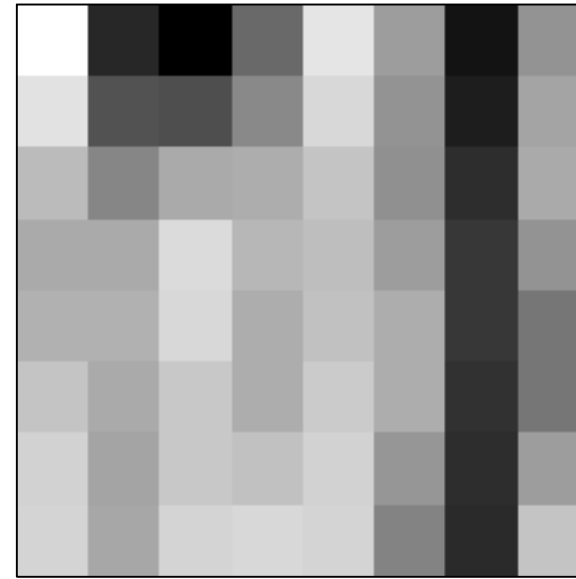
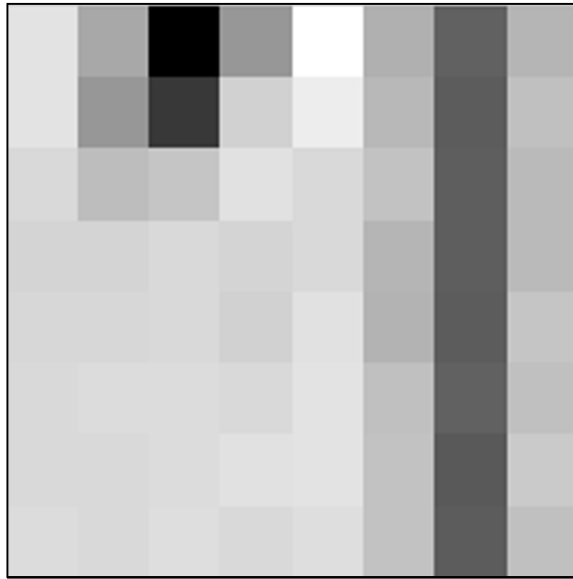
Quality Measure of a Compressed Image

Subjective quality measures (by human observer):

- This can be done by showing the compressed image to a group of viewers and averaging their evaluations. The evaluations may be made using an absolute rating scale.

Value	Rating	Description
1	Excellent	An image of extremely high quality, as good as you could desire.
2	Fine	An image of high quality, providing enjoyable viewing. Interference is not objectionable.
3	Passable	An image of acceptable quality. Interference is not objectionable.
4	Marginal	An image of poor quality; you wish you could improve it. Interference is somewhat objectionable.
5	Inferior	A very poor image, but you could watch it. Objectionable interference is definitely present.
6	Unusable	An image so bad that you could not watch it.

Original block VS. Decompressed block



- Notice the slight differences between the original (left) and decompressed image (right), which is most readily seen in the bottom-left corner.

Difference results

- The decompressed block can be compared to the original block by taking the difference results in the following error values:

-16	27	-27	0	3	-6	5	-1
-7	7	-29	13	0	0	0	-2
1	6	-2	8	-2	5	-4	-6
4	4	-9	0	0	-4	-7	1
3	3	-8	2	2	-10	-8	14
-2	7	-2	5	0	-5	-4	12
-6	8	-2	2	-2	3	-6	4
-6	7	-5	-8	-5	9	-4	-12

- With an average absolute error of about 9 values per pixels.

$$\text{(i.e., } \frac{1}{64} \sum_{x=0}^7 \sum_{y=0}^7 |e(x,y)| = 8.4705$$

JPEG compression results

- Adjust Quantization Step to Achieve Tradeoff between CR and distortion.



Original: 100KB



JPEG: 9KB



JPEG: 5KB

- Artifacts:
Inside blocks: blurring (why?); Across blocks: blocking (why?)

JPEG compression results



File size = 83,261 byte, Highest quality (Q = 100).



File size = 15,138 byte, High quality (Q = 50).

JPEG compression results



File size = 83,261 byte, Highest quality (Q = 100).



File size = 15,138 byte, High quality (Q = 50).

JPEG compression results



File size = 9,553 byte, Medium quality (Q = 25).



File size = 4,787 byte, Low quality (Q = 10).

JPEG compression results



File size = 1,523 byte, Lowest quality (Q = 1).

JPEG compression results



File size = 83,261 byte



File size = 15,138 byte



File size = 9,553 byte



File size = 4,787 byte



File size = 1,523 byte

Thanks for your attention